

# Decision Trees

**Decision tree is a type of supervised learning algorithm with a pre-defined target variable used for classification problems.**

**There are two types:**

- Categorical Variable Decision tree + Student will play baseball: Y or N
- Continuous Variable Decision Tree

## Advantages

- Easy
- Used in data exploration
- Less data cleaning required
- Data type is not a constraint
- Non parametric method

## Disadvantages

- Overfitting
- Not fit for continuous variables

## Node splits based on algorithms

1. Gini Index
  - Works well with categorical target: 1 or 0 + Binary Split
  - CART uses Gini method to create binary splits
  - For sub-nodes, uses sum of squares:  $P^2 + Q^2$
2. Chi-Square
  - Uses Chi Square
  - Generates tree call CHAID
  - To find out the statistical significance between the differences between sub-nodes and parent node
3. Information Gain
4. Reduction in Variance

## Key Parameters

1. Min\_samples\_split
  - Defines the minimum number of samples (or observations) which are required in a node to be considered for splitting.

- Used to control over-fitting. Higher values prevent a model from learning relations which might be highly specific to the particular sample selected for a tree.
  - Too high values can lead to under-fitting hence, it should be tuned using CV.
2. `Min_sample_leaf`
- Defines the minimum samples (or observations) required in a terminal node or leaf.
  - Used to control over-fitting similar to `min_samples_split`.
  - Generally lower values should be chosen for imbalanced class problems because the regions in which the minority class will be in majority will be very small
  - The maximum depth of a tree.
  - Used to control over-fitting as higher depth will allow model to learn relations very specific to a particular sample.
  - Should be tuned using CV.
4. `Max_features`
- The maximum number of terminal nodes or leaves in a tree.
  - Can be defined in place of `max_depth`. Since binary trees are created, a depth of 'n' would produce a maximum of  $2^n$  leaves.

## Need to prune tree

- We first make the decision tree to a large depth.
- Then we start at the bottom and start removing leaves which are giving us negative returns when compared from the top.
- Suppose a split is giving us a gain of say -10 (loss of 10) and then the next split on that gives us a gain of 20. A simple decision tree will stop at step 1 but in pruning, we will see that the overall gain is +10 and keep both leaves.

## Quick Examples of Decision Trees

Example 1: Titanic Data Set From Kaggle

```
#Data Set Up
train <- read.csv("train.csv", stringsAsFactors=FALSE)
test <- read.csv("test.csv", stringsAsFactors=FALSE)
attach(test)
attach(train)

## The following objects are masked from test:
##
##      Age, Cabin, Embarked, Fare, Name, Parch, PassengerId, Pclass,
##      Sex, SibSp, Ticket

#Look at data
head(train)

## PassengerId Survived Pclass
## 1             1         0      3
```

```

## 2      2      1      1
## 3      3      1      3
## 4      4      1      1
## 5      5      0      3
## 6      6      0      3
##
##                               Name      Sex Age SibSp
## 1                               Braund, Mr. Owen Harris    male  22      1
## 2 Cumings, Mrs. John Bradley (Florence Briggs Thayer) female  38      1
## 3                               Heikkinen, Miss. Laina female  26      0
## 4 Futrelle, Mrs. Jacques Heath (Lily May Peel) female  35      1
## 5                               Allen, Mr. William Henry    male  35      0
## 6                               Moran, Mr. James          male  NA      0
##  Parch      Ticket      Fare Cabin Embarked
## 1      0      A/5 21171  7.2500      S
## 2      0      PC 17599 71.2833    C85      C
## 3      0 STON/O2. 3101282  7.9250      S
## 4      0      113803 53.1000   C123      S
## 5      0      373450  8.0500      S
## 6      0      330877  8.4583      Q

str(train)

## 'data.frame':    891 obs. of  12 variables:
## $ PassengerId: int  1 2 3 4 5 6 7 8 9 10 ...
## $ Survived   : int  0 1 1 1 0 0 0 0 1 1 ...
## $ Pclass     : int  3 1 3 1 3 3 1 3 3 2 ...
## $ Name       : chr  "Braund, Mr. Owen Harris" "Cumings, Mrs. John Bradley
(Florence Briggs Thayer)" "Heikkinen, Miss. Laina" "Futrelle, Mrs. Jacques He
ath (Lily May Peel)" ...
## $ Sex       : chr  "male" "female" "female" "female" ...
## $ Age       : num  22 38 26 35 35 NA 54 2 27 14 ...
## $ SibSp     : int  1 1 0 1 0 0 0 3 0 1 ...
## $ Parch     : int  0 0 0 0 0 0 0 1 2 0 ...
## $ Ticket    : chr  "A/5 21171" "PC 17599" "STON/O2. 3101282" "113803" ..
.
## $ Fare      : num  7.25 71.28 7.92 53.1 8.05 ...
## $ Cabin     : chr  "" "C85" "" "C123" ...
## $ Embarked  : chr  "S" "C" "S" "S" ...

#Percent of those who survived
table(train$Survived)

##
##  0  1
## 549 342

prop.table(table(train$Survived)) ###38% Survived

##
##      0      1
## 0.6161616 0.3838384

```

```

#Everyone dies prediction
test$Survived <- rep(0,418)
### rep function will create the survived column in data frame and repeat
### '0' prediction 418
### very simple prediction

#Create new data frame with Pass ID and survival
submit <- data.frame(PassengerId = test$PassengerId, Survived = test$Survived
)
write.csv(submit, file = "theyallperish.csv", row.names = FALSE)

#Sex and Survival
summary(train$Sex)

##      Length      Class      Mode 
##      891 character character 

table(train$Sex)

##
## female   male
##    314    577

prop.table(table(train$Sex,train$Survived)) #divides by total num of passenge
r not within groups

##
##              0              1
## female 0.09090909 0.26150393
## male   0.52525253 0.12233446

prop.table(table(train$Sex,train$Survived),1) #denominator with gender ##most
women survived

##
##              0              1
## female 0.2579618 0.7420382
## male   0.8110919 0.1889081

##New variable for model
test$Survived <- 0
test$Survived[test$Sex == 'female'] <- 1

#Age and survival
summary(train$Age)

##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.      NA's
##      0.42  20.12   28.00   29.70   38.00   80.00     177

##Categorical assigments to Age
train$Child <- 0

```

```
train$Child[train$Age < 18] <- 1
```

*#Find the number of survivors for the Age and Gender Subsets*

```
aggregate(Survived ~ Child + Sex, data=train, FUN=sum)
```

```
##   Child   Sex Survived
## 1     0 female     195
## 2     1 female     38
## 3     0  male     86
## 4     1  male     23
```

*##For percent*

```
aggregate(Survived ~ Child + Sex, data=train, FUN=function(x) {sum(x)/length(x)})
```

```
##   Child   Sex Survived
## 1     0 female 0.7528958
## 2     1 female 0.6909091
## 3     0  male 0.1657033
## 4     1  male 0.3965517
```

*#Fare paid + Survival*

```
train$Fare2 <- '30+'
```

```
train$Fare2[train$Fare < 30 & train$Fare >= 20] <- '20-30'
```

```
train$Fare2[train$Fare < 20 & train$Fare >= 10] <- '10-20'
```

```
train$Fare2[train$Fare < 10] <- '<10'
```

*#Class/Fare/Sex affects*

```
aggregate(Survived ~ Fare2 + Pclass + Sex, data=train, FUN=function(x) {sum(x)/length(x)})
```

```
##   Fare2 Pclass   Sex Survived
## 1 20-30      1 female 0.8333333
## 2 30+       1 female 0.9772727
## 3 10-20      2 female 0.9142857
## 4 20-30      2 female 0.9000000
## 5 30+       2 female 1.0000000
## 6 <10       3 female 0.5937500
## 7 10-20      3 female 0.5813953
## 8 20-30      3 female 0.3333333
## 9 30+       3 female 0.1250000
## 10 <10      1  male 0.0000000
## 11 20-30      1  male 0.4000000
## 12 30+       1  male 0.3837209
## 13 <10       2  male 0.0000000
## 14 10-20      2  male 0.1587302
## 15 20-30      2  male 0.1600000
## 16 30+       2  male 0.2142857
## 17 <10       3  male 0.1115385
## 18 10-20      3  male 0.2368421
```


```
## 19 20-30      3   male 0.1250000
## 20  30+      3   male 0.2400000

## Why would someone in third class with an expensive ticket be worse off?
test$Survived <- 0
test$Survived[test$Sex == 'female'] <- 1
test$Survived[test$Sex == 'female' & test$Pclass == 3 & test$Fare >= 20] <- 0

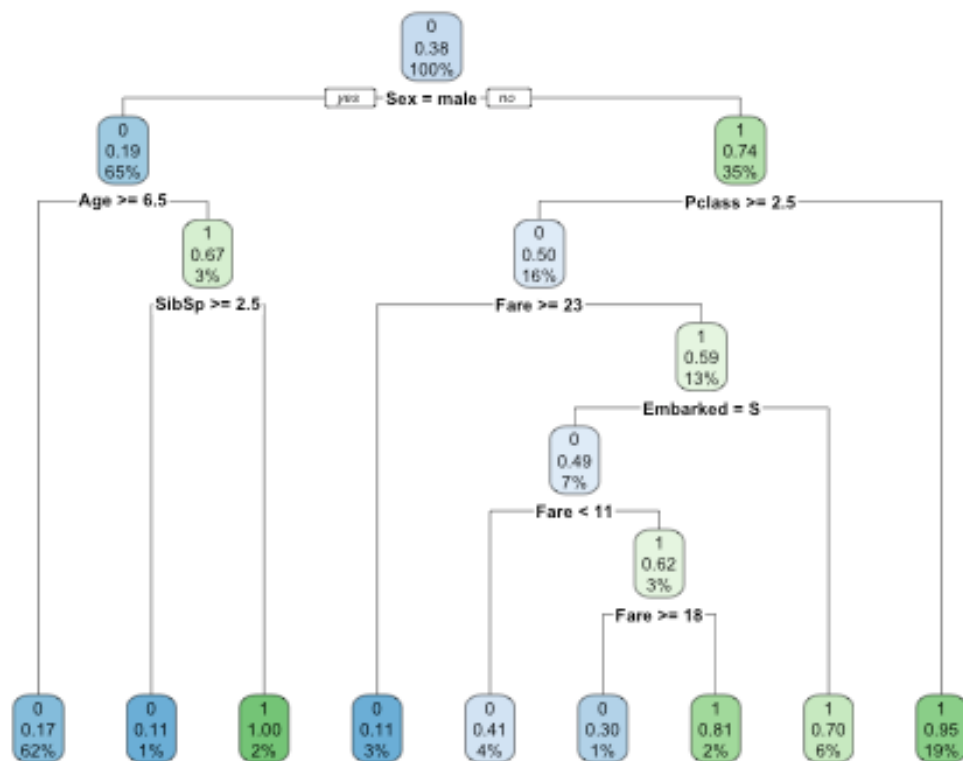
#Recursive Partitioning and Regression Trees
library(rpart)

## Warning: package 'rpart' was built under R version 3.4.3

## to predict a continuous variable use method = "anova"
fit <- rpart(Survived ~ Pclass + Sex + Age + SibSp + Parch + Fare + Embarked,
data=train,method="class")
plot(fit)
text(fit)
```



```
#For prettier tree
library(rpart.plot)
library(RColorBrewer)
rpart.plot(fit)
```



Random Forest gives more accurate predictions.